



NVIDIA Training Course Catalog

November 2023



Introduction

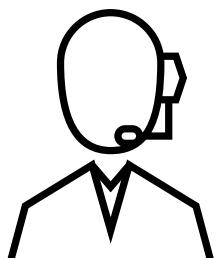
NVIDIA offers training for diverse needs, giving individuals and teams across organizations what they need to advance their knowledge in AI, accelerated computing, data science, data center administration, graphics and simulation, networking, and more.

With access to high-performance computing, you'll learn how to train, optimize, and deploy neural networks using the latest deep learning tools, frameworks, and SDKs. You'll also learn how to assess, parallelize, optimize, and deploy GPU-accelerated computing applications.

Our training program offers both self-paced online courses and instructor-led, prescheduled workshops. The self-paced courses range from 10 minutes to 8 hours and guide you through applying a specific technology, setting up a project, or administering solutions in a data center. Instructor-led workshops and boot camps go deeper into topic areas, teaching you how to implement a project or solution from end to end. Both types of courses give you valuable hands-on experience using the latest technologies.

Why Choose NVIDIA for Training?

- Learn how to build deep learning and accelerated computing applications for industries such as healthcare, robotics, autonomous driving, manufacturing, and more.
- Gain hands-on experience with the most widely used, industry-standard platforms including software, hardware, tools, and frameworks. Each student will have access to a fully configured, GPU-accelerated server in the cloud or access to NVIDIA solutions in our training lab.
- Become proficient in administering NVIDIA hardware and software solutions such as DGX™, InfiniBand, Cumulus, NVIDIA AI Enterprise, and more.
- Access instructor-led workshops and online courses from anywhere using just a laptop and internet connection.
- Acquire real-world expertise through content designed in collaboration with industry leaders such as Children's Hospital of Los Angeles, Mayo Clinic, and PwC.
- Earn NVIDIA certifications and course completion certificates to indicate subject matter competency and support your career growth.



For team training, contact an **NVIDIA training advisor**, who will work with you to create a customized plan that addresses your team's specific training needs and is aligned to your business objectives and priorities.

Table of Contents

Instructor-Led Workshops for Developers

Accelerated Computing

Accelerating CUDA® C++ Applications With Multiple GPUs	7
Fundamentals of Accelerated Computing With CUDA C/C++	7
Fundamentals of Accelerated Computing With CUDA Python	7
Fundamentals of Accelerated Computing With OpenACC®	7
Scaling CUDA C++ Applications to Multiple Nodes	8

Data Science

Accelerating Data Engineering Pipelines	8
Enhancing Data Science Outcomes With Efficient Workflows	8
Fundamentals of Accelerated Data Science	8

Deep Learning

Applications of AI for Anomaly Detection	9
Applications of AI for Predictive Maintenance	9
Building AI-Based Cybersecurity Pipelines	9
Building Conversational AI Applications V2.0	10
Building Deep Learning-Based Anti-Fraud Applications	10
Building Transformer-Based Natural Language Processing	10
Computer Vision for Industrial Inspection	10
Data Parallelism: How to Train Deep Learning Models on Multiple GPUs	11
Fundamentals of Deep Learning	11
Model Parallelism: Building and Deploying Large Neural Networks	11

Generative AI and Large Language Models (LLMs)

Efficient Large Language Model Customizations	12
Generative AI With Diffusion Models	12
Rapid Application Development Using Large Language Models	12

Graphics and Simulation

Bootstrapping Computer Vision Models with Synthetic Data	13
Building Digital Avatar Pipelines With NVIDIA Omniverse Audio2Face and Riva	13

Online, Self-Paced Courses for Developers

Accelerated Computing Fundamentals

Accelerating CUDA C++ Applications With Concurrent Streams	13
An Even Easier Introduction to CUDA	13
Fundamentals of Accelerated Computing With CUDA Python	14
Fundamentals of Accelerated Computing With OpenACC	14

Getting Started With Accelerated Computing With CUDA C/C++	14
GPU Acceleration With the C++ Standard Library	14
High-Performance Computing With Containers	14
Optimizing CUDA Machine Learning Codes With NVIDIA Nsight™ Profiling Tools	15
Scaling GPU-Accelerated Applications With the C++ Standard Library	15
Scaling Workloads Across Multiple GPUs With CUDA C++	15
 Data Science	
Accelerating End-to-End Data Science Workflows	16
RAPIDS Accelerator for Apache Spark	16
 Deep Learning	
Building a Brain in 10 Minutes	16
Building Real-Time Video AI Applications	16
Building Video AI Applications at the Edge on NVIDIA® Jetson Nano™	16
Deploying a Model for Inference at Production Scale	17
Digital Fingerprinting With Morpheus	17
Disaster Risk Monitoring Using Satellite Imagery	17
Get Started With Highly Accurate Custom ASR for Speech AI	17
Getting Started With AI on Jetson Nano	17
Getting Started With Deep Learning	18
Getting Started With Image Segmentation	18
Integrating Sensors With NVIDIA DRIVE	18
Introduction to Graph Neural Networks	18
Introduction to Physics-Informed Machine Learning With NVIDIA Modulus	18
Modeling Time-Series Data With Recurrent Neural Networks in Keras	19
 Generative AI and Large Language Models (LLMs)	
Generative AI Explained	19
Generative AI With Diffusion Models	19
Introduction to Transformer-Based Natural Language Processing	19
Prompt Engineering With Llama 2	19
 Graphics and Simulation	
Assemble a Simple Robot in NVIDIA Isaac Sim™	20
Build Beautiful, Custom UI for 3D Tools on NVIDIA Omniverse™	20
Develop, Customize, and Publish in NVIDIA Omniverse With Extensions	20
Easily Develop Advanced 3D Layout Tools on NVIDIA Omniverse	20
Essentials of Developing Omniverse Kit Applications	21
Getting Started With USD for Collaborative 3D Workflows	21
How to Build Customer 3D Scene Manipulator Tools on NVIDIA Omniverse	21

Introduction to Robotic Simulations in NVIDIA Isaac Sim	22
Synthetic Data Generation for Training Computer Vision Models	22

Infrastructure

Introduction to AI in the Data Center	23
Introduction to NVIDIA DOCA™ for DPUs	23
Getting Started With DOCA Flow	23

Instructor-Led Workshops for Administrators

AI and Data Science

NVIDIA AI Enterprise Administration: Public Training	24
--	----

Cluster Administration

Base Command Manager	24
----------------------	----

Ethernet Cumulus

Cumulus® Linux: Public Bootcamp	24
Cumulus Linux: Private Workshop	24
NVIDIA Cumulus Linux: Customized Advanced Training	24

InfiniBand

InfiniBand Customized Course	25
Networking Professional Customized Training	25

NVIDIA DGX

NVIDIA DGX H100/A100 Administration: Private Workshop	25
NVIDIA DGX H100/A100 Administration: Public Workshop	25
NVIDIA DGX BasePOD Administration: Private Workshop	26
NVIDIA DGX SuperPOD™ Administration: Private Workshop	26

Virtualization

NVIDIA AI Enterprise Administration: Public Bootcamp	26
--	----

Online, Self-Paced Courses for Administrators

AI and Data Science

Introduction to AI in the Data Center	27
NVIDIA AI Enterprise Administration	27
NVIDIA NeMo Framework	27

Cluster Administration

Base Command Manager	27
Base Command Manager Autoscaling Hybrid Cloud	28
Introduction to Base Command Manager	28

DGX	
NVIDIA DGX Cloud	28
Ethernet	
Linux Networking Fundamentals	28
Network Administration With the NVIDIA Onyx™ Switch System	29
RDMA Over Converged Ethernet (RoCE) From A to Z	29
Graphics and Simulation	
NVIDIA Omniverse Enterprise Administration	29
InfiniBand	
InfiniBand Essentials	29
InfiniBand Professional	30
Management	
Data Center Management Made Easy With NVIDIA UFM	30
NVIDIA License System	30
Network	
Ansible Essentials for Network Engineers	30
Introduction to Networking	31
MLXlink and MLXcables Debug Tools	31
NVIDIA BlueField DPU Administration	31
RDMA	
The Fundamentals of RDMA Programming	31
Certifications	
NVIDIA Certified Associate: AI in the Data Center	32
NVIDIA Certified Professional: InfiniBand	32

Instructor-Led Workshops for Developers

Workshop Name	Description	Prerequisites			
Accelerated Computing					
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
Accelerating CUDA® C++ Applications With Multiple GPUs	Discover how to write CUDA C++ applications that efficiently and correctly use all available GPUs in a single node, dramatically improving the performance of applications and making the most cost-effective use of systems with multiple GPUs.	Professional experience programming CUDA C/C++ applications, including the use of the NVIDIA CUDA Compiler (NVCC), kernel launches, grid-stride loops, host-to-device and device-to-host memory transfers, and CUDA error handling. Familiarity with the Linux command line and experience using makefiles to compile C/C++ code.			
	> Learn More	CUDA C++, NVCC, Nsight Systems	English, Simplified Chinese	8 hours \$500 (excludes tax, if applicable)	Yes
Fundamentals of Accelerated Computing With CUDA C/C++	Learn how to accelerate and optimize existing C/C++ CPU-only applications to apply the power of GPUs using the most essential CUDA techniques and the NVIDIA Nsight Systems profiler.	Basic C/C++ competency, including familiarity with variable types, loops, conditional statements, functions, and array manipulations. No previous knowledge of CUDA programming is assumed.			
	> Learn More	NVIDIA Nsight Systems, nsys	English, Korean, Japanese, Simplified Chinese, Traditional Chinese	8 hours \$500 (excludes tax, if applicable)	Yes
Fundamentals of Accelerated Computing With CUDA Python	Explore how to use Numba—the just-in-time, type-specializing Python function compiler—to create and launch CUDA kernels to accelerate Python programs on massively parallel NVIDIA GPUs.	Basic Python competency, including familiarity with variable types, loops, conditional statements, functions, and array manipulations. Also, must have NumPy competency, including the use of ndarrays and ufuncs.			
	> Learn More	CUDA, Python, Numba, NumPy	English, Simplified Chinese, Traditional Chinese	8 hours \$500 (excludes tax, if applicable)	Yes
Fundamentals of Accelerated Computing With OpenACC®	Find out how to write and configure code parallelization with OpenACC, optimize memory movements between the CPU and GPU accelerator, and apply the techniques to accelerate a CPU-only Laplace heat equation to achieve performance gains.	Basic C/C++ or Fortran competency, including familiarity with variable types, loops, conditional statements, functions, and array manipulations. No previous knowledge of GPU programming is assumed.			
	> Learn More	NVIDIA Nsight, OpenACC	English	8 hours \$500 (excludes tax, if applicable)	Yes

[Back](#)

Workshop Name	Description	Prerequisites				
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
Scaling CUDA C++ Applications to Multiple Nodes	Learn the tools and techniques needed to write CUDA C++ applications that can scale efficiently to clusters of NVIDIA GPUs.		English, Simplified Chinese	8 hours	\$500 (excludes tax, if applicable)	Yes

Data Science

Accelerating Data Engineering Pipelines	Explore how to employ advanced data engineering tools and techniques with GPUs to significantly improve data engineering pipelines.	Intermediate knowledge of Python (list comprehension, objects). Familiarity with pandas and introductory statistics (mean, median, mode) a plus.
Enhancing Data Science Outcomes With Efficient Workflows	Learn how to create an end-to-end, hardware-accelerated machine learning pipeline for large datasets. Throughout the development process, you'll use diagnostic tools to identify delays and learn to mitigate common pitfalls.	<ul style="list-style-type: none"> > Basic knowledge of a standard data science workflow on tabular data. > Knowledge of distributed computing using Dask. > Completion of the DLI's Fundamentals of Accelerated Data Science course or an ability to manipulate data using cuDF and some experience building machine learning models using cuML.
Fundamentals of Accelerated Data Science	Learn how to perform multiple analysis tasks on large datasets using NVIDIA RAPIDS™, a collection of data science libraries that allows end-to-end GPU acceleration for data science workflows.	Professional data science experience with Python, including proficiency in pandas and NumPy. Also, must have familiarity with common machine learning algorithms, including XGBoost, linear regression, DBSCAN, K-Means, and SSSP.

[Back](#)

Workshop Name	Description	Prerequisites				
Deep Learning						
Workshop Name	Description	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
Applications of AI for Anomaly Detection	Learn to detect anomalies in large datasets to identify network intrusions using supervised and unsupervised machine learning techniques, such as accelerated XGBoost, autoencoders, and generative adversarial networks (GANs).	RAPIDS, XGBoost, TensorFlow, Keras, pandas, autoencoders, GANs	English	8 hours	\$500 (excludes tax, if applicable)	Yes
	> Learn More					
Applications of AI for Predictive Maintenance	Discover how to identify anomalies and failures in time-series data, estimate the remaining useful life of the corresponding parts, and use this information to map anomalies to failure conditions.	Python, TensorFlow, Keras, XGBoost, RAPIDS, cuDF, long short-term memory (LSTM), autoencoders	English, Simplified Chinese	8 hours	\$500 (excludes tax, if applicable)	Yes
	> Learn More					
Building AI-Based Cybersecurity Pipelines	Traditional cybersecurity methods include creating barriers around your infrastructure to protect it from intruders. However, as enterprises continue to digitally transform, they're faced with a proliferation of devices, more sophisticated cybersecurity attacks, and an incredibly vast network of data to protect—which means new cybersecurity methodologies must be explored. An alternative approach is to address cybersecurity as a data science problem: Better understand all the users and activities across your network so that you can identify which transactions are typical and which are potentially nefarious. The NVIDIA Morpheus AI framework lets cybersecurity developers and practitioners harness the power of GPU computing to implement cybersecurity solutions that perform on a scale never before possible. With Morpheus, cybersecurity developers can create optimized AI pipelines for filtering, processing, and classifying large volumes of real-time data. Bringing a new level of information security to data centers, Morpheus enables dynamic protection, real-time telemetry, and adaptive defenses for detecting and remediating cybersecurity threats.					
	> Learn More					
Workshop Name	Description	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
		NVIDIA Morpheus, NVIDIA Triton Inference Server, RAPIDS, CLX, Helm, Kubernetes	English	8 hours	\$500 (excludes tax, if applicable)	Yes

Back

Workshop Name	Description	Prerequisites			
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
Building Conversational AI Applications V2.0	Discover how to quickly build and deploy production-quality speech AI applications with real-time transcription and natural language processing capabilities. > Learn More	NVIDIA Riva, NVIDIA TAO Toolkit, Kubernetes	English	8 hours	\$500 (excludes tax, if applicable)
Building Deep Learning-Based Anti-Fraud Applications	This course is primarily for data scientists and professionals working in the field of financial fraud modeling in banks. It teaches how to train, accelerate, and optimize fraud detection classifiers based on machine learning and deep learning. > Learn More				<ul style="list-style-type: none"> > Basic Python programming experience. > Fundamental understanding of deep learning frameworks (such as TensorFlow, PyTorch, or Keras). > Basic knowledge of neural networks.
Building Transformer-Based Natural Language Processing	RAPIDS, CuPy, PyTorch, Deep Graph Library, NVIDIA NeMo™, NVIDIA Triton Inference Server	Simplified Chinese	8 hours	\$500 (excludes tax, if applicable)	Yes
Computer Vision for Industrial Inspection	In this workshop, you'll learn how Transformers are used as the building blocks of modern large language models (LLMs). You'll then use these models for various NLP tasks, including text classification, named-entity recognition (NER), author attribution, and question answering. You'll also learn how to analyze various model features, constraints, and characteristics to determine which model is best suited for a particular use case based on metrics, domain specificity, and available resources. > Learn More	PyTorch, pandas, NVIDIA NeMo, NVIDIA Triton Inference Server	English, Simplified Chinese	8 hours	\$500 (excludes tax, if applicable)
	Experience with Python coding and use of library functions and parameters. Fundamental understanding of a deep learning framework, such as TensorFlow, PyTorch, or Keras. And basic understanding of neural networks.				
	<ul style="list-style-type: none"> > Experience with Python; basic understanding of data processing and deep learning > To gain experience with Python, we suggest this Python tutorial > For a basic understanding of data processing and deep learning, we suggest Fundamentals of Deep Learning. 				
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	Python, pandas, DALI, NVIDIA TAO Toolkit, NVIDIA TensorRT™, and NVIDIA Triton Inference Server	English, Simplified Chinese	8 hours	\$500 (excludes tax, if applicable)	Yes

[Back](#)

Workshop Name	Description	Prerequisites			
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
Data Parallelism: How to Train Deep Learning Models on Multiple GPUs	This workshop teaches you techniques for data-parallel deep learning training on multiple GPUs to shorten the training time required for data-intensive applications. Working with deep learning tools, frameworks, and workflows to perform neural network training, you'll learn how to decrease model training time by distributing data to multiple GPUs, while retaining the accuracy of training on a single GPU.	English, Simplified Chinese	8 hours	\$500 (excludes tax, if applicable)	Yes
> Learn More					
Fundamentals of Deep Learning	Learn how deep learning works through hands-on exercises in computer vision and natural language processing (NLP). You'll train deep learning models from scratch and pick up tricks and tools for achieving highly accurate results along the way. You'll also learn to leverage freely available, state-of-the-art pretrained models to save time and get your deep learning application up and running quickly.	Tensorflow, Keras, pandas, NumPy	English, Simplified Chinese, Japanese	8 hours	\$500 (excludes tax, if applicable)
> Learn More					
Model Parallelism: Building and Deploying Large Neural Networks	In this workshop, you'll learn how to scale training and deployment of LLMs and neural networks across multiple nodes, use various forms of model parallelism to overcome the challenges associated with large-model memory footprint, capture and understand training performance characteristics to optimize model architecture and deploy very large multi-GPU, multi-node models to production using NVIDIA Triton™ Inference Server.	PyTorch, Megatron-LM, DeepSpeed, Slurm, NVIDIA Triton Inference Server, NVIDIA Nsight	English, Korean, Simplified Chinese	8 hours	\$500 (excludes tax, if applicable)
> Learn More					

Back

Generative AI and Large Language Models (LLMs)

Efficient Large Language Model Customizations

Learn a variety of techniques to efficiently customize pretrained LLMs for your specific use cases—without engaging in the computationally intensive and expensive process of pretraining your own model or fine-tuning a model's internal weights. Using the open-source NVIDIA NeMo framework, you'll learn prompt engineering and various parameter-efficient fine-tuning methods to customize LLM behavior for your organization.

[> Learn More](#)

- > Professional experience with the Python programming language.
- > Familiarity with fundamental deep learning topics like model architecture, training and inference.
- > Familiarity with a modern Python-based deep learning framework (PyTorch preferred).
- > Familiarity working with out-of-the-box pretrained LLMs.

Tools, Libraries, Frameworks

Python, NVIDIA NeMo, GPT, LLaMA, HuggingFace

Languages

English

Duration

8 hours

Price

\$500 (excludes tax, if applicable)

Certificate

Yes

Generative AI With Diffusion Models

Get started with gen AI application development with this hands-on course where you'll learn how to build a text-to-image generative AI application using the latest techniques. Generate images with diffusion models and refine the output with various optimizations. Build a denoising diffusion model from the U-Net architecture to context embeddings for greater user control.

[> Learn More](#)

- > Good understanding of **PyTorch**
- > Good understanding of **deep learning**

Tools, Libraries, Frameworks

PyTorch, CLIP

Languages

English

Duration

8 hours

Price

\$500 (excludes tax, if applicable)

Certificate

Yes

Rapid Application Development Using Large Language Models

In this course, you'll gain a strong understanding and practical knowledge of LLM application development by exploring the open-source ecosystem, including pretrained LLMs, that can help you get started quickly developing LLM-based applications.

[> Learn More](#)

- > Introductory deep learning, with comfort with PyTorch and transfer learning preferred. Content covered by DL's Getting Started With Deep Learning or Fundamentals of Deep Learning courses, or similar experience is sufficient.
- > Intermediate Python experience, including object-oriented programming and libraries. Content covered by Python Tutorial ([w3schools.com](https://www.w3schools.com/python/)) or similar experience is sufficient.

Tools, Libraries, Frameworks

Python, PyTorch, HuggingFace, transformers, LangChain, LlamaIndex

Languages

English

Duration

8 hours

Price

\$500 (excludes tax, if applicable)

Certificate

Yes

[Back](#)

Graphics and Simulation

Bootstrapping Computer Vision Models with Synthetic Data	Learn how to use NVIDIA Omniverse Replicator, a core Omniverse extension, to accelerate the development of computer vision models. Generate accurate, photorealistic, physics-conforming synthetic data to ease the expensive, time-consuming task of labeling real-world data. Omniverse Replicator accelerates AI development at scale and reduces time to production.	> Learn More	> Intermediate understanding of Python (including classes, objects, and decorators). ► Basic understanding of Machine Learning and Deep Learning concepts and pipelines.										
	<table border="1"> <thead> <tr> <th>Tools, Libraries, Frameworks</th><th>Languages</th><th>Duration</th><th>Price</th><th>Certificate</th></tr> </thead> <tbody> <tr> <td>Omniverse Replicator, Omniverse Defect Extension</td><td>English</td><td>8 hours</td><td>\$500 (excludes tax, if applicable)</td><td>Yes</td></tr> </tbody> </table>	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate	Omniverse Replicator, Omniverse Defect Extension	English	8 hours	\$500 (excludes tax, if applicable)	Yes		
Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate									
Omniverse Replicator, Omniverse Defect Extension	English	8 hours	\$500 (excludes tax, if applicable)	Yes									
Building Digital Avatar Pipelines With NVIDIA Omniverse Audio2Face and Riva	This course, from an end-to-end application development perspective, will provide you with detailed guidance on how to use NVIDIA Omniverse Audio2Face and the interactive speech suite Riva to build virtual digital humans.	> Learn More	> Basic Python programming experience. ► Fundamental understanding of deep neural networks.										
	<table border="1"> <thead> <tr> <th>Tools, Libraries, Frameworks</th><th>Languages</th><th>Duration</th><th>Price</th><th>Certificate</th></tr> </thead> <tbody> <tr> <td>NVIDIA Omniverse Audio2Face, NVIDIA Riva, PyTorch</td><td>Simplified Chinese</td><td>8 hours</td><td>\$500 (excludes tax, if applicable)</td><td>Yes</td></tr> </tbody> </table>	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate	NVIDIA Omniverse Audio2Face, NVIDIA Riva, PyTorch	Simplified Chinese	8 hours	\$500 (excludes tax, if applicable)	Yes		
Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate									
NVIDIA Omniverse Audio2Face, NVIDIA Riva, PyTorch	Simplified Chinese	8 hours	\$500 (excludes tax, if applicable)	Yes									

Online, Self-Paced Courses for Developers

Course Name	Description	Prerequisites										
Accelerated Computing Fundamentals												
Accelerating CUDA C++ Applications With Concurrent Streams	Discover how to improve performance for your CUDA C/C++ applications by overlapping memory transfers to and from the GPU with computations on the GPU.	> Learn More										
		<table border="1"> <thead> <tr> <th>Tools, Libraries, Frameworks</th><th>Languages</th><th>Duration</th><th>Price</th><th>Certificate</th></tr> </thead> <tbody> <tr> <td></td><td>English</td><td>4 hours</td><td>\$30 (excludes tax, if applicable)</td><td>Yes</td></tr> </tbody> </table>	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate		English	4 hours	\$30 (excludes tax, if applicable)	Yes
Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate								
	English	4 hours	\$30 (excludes tax, if applicable)	Yes								
An Even Easier Introduction to CUDA	Learn the basics of writing parallel CUDA kernels to run on NVIDIA GPUs.	> Learn More										
		<table border="1"> <thead> <tr> <th>Tools, Libraries, Frameworks</th><th>Languages</th><th>Duration</th><th>Price</th><th>Certificate</th></tr> </thead> <tbody> <tr> <td>C/C++</td><td>English</td><td>1 hour</td><td>Free</td><td>N/A</td></tr> </tbody> </table>	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate	C/C++	English	1 hour	Free	N/A
Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate								
C/C++	English	1 hour	Free	N/A								

[Back](#)

Course Name	Description	Prerequisites			
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
Fundamentals of Accelerated Computing With CUDA Python	Explore how to use Numba—the just-in-time, type-specializing Python function compiler—to create and launch CUDA kernels to accelerate Python programs on massively parallel NVIDIA GPUs.	Basic Python competency, including familiarity with variable types, loops, conditional statements, functions, and array manipulations. Also, must have NumPy competency, including the use of ndarrays and ufuncs.			
	> Learn More				
Fundamentals of Accelerated Computing With OpenACC	Find out how to build and optimize accelerated heterogeneous applications on multiple GPU clusters using a combination of OpenACC, CUDA-aware MPI, and NVIDIA profiling tools.				Basic experience with C/C++
	> Learn More				
Getting Started With Accelerated Computing With CUDA C/C++	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	OpenACC, C/C++	English	8 hours	\$90 (excludes tax, if applicable)	N/A
GPU Acceleration With the C++ Standard Library	Discover how to accelerate and optimize existing C/C++ CPU-only applications to leverage the power of GPUs using the most essential CUDA techniques and the Nsight Systems profiler.				Basic C/C++ competency, including familiarity with variable types, loops, conditional statements, functions, and array manipulations. No previous knowledge of CUDA programming is assumed.
	> Learn More				
High-Performance Computing With Containers	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	C/C++, CUDA	English, Japanese, Korean, Simplified Chinese, Traditional Chinese	8 hours	\$90 (excludes tax, if applicable)	Yes
GPU Acceleration With the C++ Standard Library	Learn to write simple, portable, parallel-first applications using only standard C++ language features that can be compiled without modification to take advantage of NVIDIA GPU-accelerated environments.				Beginner-level experience with C++11 . Comfort working with C++ lambdas and standard library algorithms .
	> Learn More				
High-Performance Computing With Containers	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	C++, NVIDIA HPC SDK	English	2 hours	\$30 (excludes tax, if applicable)	N/A
High-Performance Computing With Containers	Learn how to reduce complexity and improve portability and efficiency of your code by using a containerized environment for HPC application development.				Proficiency in programming in C/C++ and professional experience working on HPC applications.
	> Learn More				
High-Performance Computing With Containers	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	Docker, Singularity, HPC Container Maker (HPCCM), C/C++	English	2 hours	\$30 (excludes tax, if applicable)	N/A

Back

Course Name	Description	Prerequisites		
Optimizing CUDA Machine Learning Codes With NVIDIA Nsight™ Profiling Tools	<p>NVIDIA Developer Tools are a collection of applications, spanning desktop and mobile targets, that enable developers to build, debug, profile, and develop class-leading and cutting-edge software using the latest visual computing hardware from NVIDIA. In this course, you'll learn the effective use of two powerful NVIDIA developer tools: Nsight Systems and Nsight Compute.</p> <p>Nsight Systems provide developers with a system-wide visualization of an application's performance. Developers can optimize bottlenecks to scale efficiently across any number or size of CPU and GPU—from large servers to the smallest systems on chip. Nsight Compute is an interactive kernel profiler for CUDA applications. It provides detailed performance metrics and API debugging via a user interface and command-line tool.</p> <p>By the time you complete this course, you'll be able to use Nsight Systems and Nsight Compute to analyze and optimize CUDA applications. Following best practices, you'll begin by using Nsight Systems to analyze overall application structure and explore parallelization opportunities before turning to Nsight Compute to analyze and optimize individual CUDA kernels.</p>	Familiarity with machine learning applications using CUDA. We suggest Fundamentals of Accelerated Computing with CUDA C/C++ .		
Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
NVIDIA Nsight Systems , NVIDIA Nsight Compute	English	2 hours	\$30 (excludes tax, if applicable)	N/A
Scaling GPU-Accelerated Applications With the C++ Standard Library				
In this interactive, hands-on workshop, which is the followup to GPU Acceleration With the C++ Standard Library, you'll learn how to write scalable, GPU-accelerated, hybrid applications using C++ standard language features alongside MPI.	Beginner-level experience with C++ 11; comfort working with C++ lambdas and standard library algorithms; experience developing C++/MPI hybrid applications that require inter-rank communication; comfort working with C++ concurrency primitives such as std::thread, std::barrier, and std::atomic.			
Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
C++, NVIDIA HPC SDK, MPI	English	2 hours	\$30 (excludes tax, if applicable)	N/A
Scaling Workloads Across Multiple GPUs With CUDA C++			Competency writing applications in CUDA C/C++.	
Learn how to build robust and efficient CUDA C++ applications that can take advantage of all available GPUs on a single node.				
Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
C/C++, accelerated computing, CUDA	English	4 hours	\$30 (excludes tax, if applicable)	Yes

Back

Course Name	Description	Prerequisites			
Data Science					
Accelerating End-to-End Data Science Workflows	Explore how to perform multiple analysis tasks on large datasets using RAPIDS, a collection of data science libraries that allows end-to-end GPU acceleration for data science workflows.	Experience with Python, ideally including pandas and NumPy.			
	> Learn More	Tools, Libraries, Frameworks	Languages	Duration	Price
	RAPIDS, cuDF, cuML, cuGraph, Apache Arrow	English, Simplified Chinese	6 hours	\$90 (excludes tax, if applicable)	Yes
RAPIDS Accelerator for Apache Spark	In this training lab, we'll walk through the RAPIDS Accelerator for Apache Spark, including running SQL queries on CPU and GPU in Spark and diving into the toolset that helps enable success.	<ul style="list-style-type: none"> ➤ Basic experience with Linux terminal commands. ➤ Basic experience with Python ➤ Basic experience with Spark, PySpark, or pandas 			
	> Learn More	Tools, Libraries, Frameworks	Languages	Duration	Price
	RAPIDS, Spark	English	2 hours	\$30 (excludes tax, if applicable)	N/A
Deep Learning					
Building a Brain in 10 Minutes	This one-click notebook explores the biological and psychological inspirations for the world's first neural networks.	An understanding of fundamental programming concepts in Python 3 such as functions, loops, dictionaries, and arrays.			
	> Learn More	Tools, Libraries, Frameworks	Languages	Duration	Price
	N/A	English	10 minutes	Free	N/A
Building Real-Time Video AI Applications	Gain the knowledge and skills needed to enable the real-time transformation of raw video data from widely deployed camera sensors into deep learning-based insights.	Competency in the Python 3 programming language, some experience manipulating data using pandas DataFrames, and familiarity with deep networks (specifically variations of CNNs).			
	> Learn More	Tools, Libraries, Frameworks	Languages	Duration	Price
	NVIDIA DeepStream, NVIDIA TAO Toolkit, and NVIDIA TensorRT	English, Simplified Chinese	8 hours	\$90.00 (excludes tax, if applicable)	N/A
Building Video AI Applications at the Edge on NVIDIA® Jetson Nano™	Use JupyterLab notebooks and Python application samples on your Jetson Nano to build new projects that extract meaningful insights from video streams through deep learning video analytics.	Basic familiarity with the Linux command line and an understanding of fundamental programming concepts in Python 3, such as functions, loops, dictionaries, and arrays.			
	> Learn More	Tools, Libraries, Frameworks	Languages	Duration	Price
	DeepStream, TensorRT, Jetson Nano, Python	English, Simplified Chinese	8 hours	Free (hardware required)	N/A

Back

Course Name	Description	Prerequisites			
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
Deploying a Model for Inference at Production Scale	<p>Learn how to deploy your own machine learning models on a GPU server.</p> <p>> Learn More</p>	NVIDIA Triton	English	4 hours	\$30 (excludes tax, if applicable)
Digital Fingerprinting With Morpheus	<p>In this course, you'll get hands-on experience developing and deploying the NVIDIA digital fingerprinting AI workflow that enables 100% data visibility and drastically reduces the time to detect threats. You'll also hear from cybersecurity experts from a variety of institutions about how to use NVIDIA AI frameworks and tools to architect cybersecurity solutions.</p> <p>> Learn More</p>	NVIDIA Morpheus AI framework, NVIDIA Triton Inference Server	English	1 hour	Free
Disaster Risk Monitoring Using Satellite Imagery	<p>Learn how to build and deploy a deep learning model to automate the detection of flood events using satellite imagery. This workflow can be applied to lower the cost, improve the efficiency, and significantly enhance the effectiveness of various natural disaster management use cases.</p> <p>> Learn More</p>	NVIDIA DALI®, the NVIDIA TAO Toolkit, NVIDIA TensorRT, NVIDIA Triton Inference Server	English, Simplified Chinese	10 hours	Free
Get Started With Highly Accurate Custom ASR for Speech AI	<p>Learn to build, train, fine-tune, and deploy a GPU-accelerated automatic speech recognition service with NVIDIA Riva that includes customized features.</p> <p>> Learn More</p>	Riva, TAO Toolkit, Kubernetes	English	2 hours	\$30 (excludes tax, if applicable)
Getting Started With AI on Jetson Nano	<p>Discover how to build a deep learning classification project with computer vision models using the NVIDIA Jetson Nano Developer Kit.</p> <p>> Learn More</p>	PyTorch, Jetson Nano	English, Simplified Chinese, Japanese, Korean	8 hours	Free (hardware required)

Back

Course Name	Description	Prerequisites			
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
Getting Started With Deep Learning	Explore the fundamentals of deep learning by training neural networks and using results to improve performance and capabilities. > Learn More	TensorFlow 2 with Keras, pandas	English, Simplified Chinese	8 hours	\$90 (excludes tax, if applicable)
Getting Started With Image Segmentation				Basic experience training neural networks.	
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	TensorFlow 2 with Keras	English	2 hours	\$30 (excludes tax, if applicable)	N/A
Integrating Sensors With NVIDIA DRIVE				Basic experience in C++ and Linux terminal commands.	
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	C++, NVIDIA DriveWorks	English	2 hours	\$30 (excludes tax, if applicable)	N/A
Introduction to Graph Neural Networks				Competency in the Python 3 programming language. Experience with deep neural networks (specifically variations of CNNs).	
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	Deep Graph Library, PyTorch	English	2 hours	\$30 (excludes tax, if applicable)	N/A
Introduction to Physics-Informed Machine Learning With NVIDIA Modulus				<ul style="list-style-type: none"> > Familiarity with the Python programming language > An understanding of partial differential equations and their use in physics. > Familiarity with machine learning concepts like training and inference. 	
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	NVIDIA Modulus	English	4 hours	\$30 (excludes tax, if applicable)	N/A

Back

Course Name	Description	Prerequisites				
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
Modeling Time-Series Data With Recurrent Neural Networks in Keras	Explore how to classify and forecast time-series data using recurrent neural networks (RNNs), such as modeling a patient's health over time. > Learn More	Keras	English	2 hours	\$30 (excludes tax, if applicable)	N/A

Generative AI and Large Language Models (LLMs)

Generative AI Explained	Generative AI describes technologies that are used to generate new content based on a variety of inputs. In this course, you will learn Generative AI concepts, applications, as well as the challenges and opportunities in this exciting field. > Learn More	Basic understanding of Machine Learning and Deep Learning concepts
		Tools, Libraries, Frameworks Languages Duration Price Certificate
	N/A English 2 hours Free N/A	
Generative AI With Diffusion Models	In this workshop, you'll train deep learning models from scratch and learn tools and tricks to achieve highly accurate results. You'll also learn to leverage freely available, state-of-the-art pretrained models to save time and get your deep learning application up and running quickly. > Learn More	An understanding of fundamental programming concepts in Python such as functions, loops, dictionaries, and arrays.
		Tools, Libraries, Frameworks Languages Duration Price Certificate
	TensorFlow 2 with Keras, pandas English 8 hours \$90 (excludes tax, if applicable) Yes	
Introduction to Transformer-Based Natural Language Processing	In this course, you'll learn how transformers are used as the building blocks of modern large language models (LLMs). You'll then use these models for various NLP tasks, including text classification, named-entity recognition (NER), author attribution, and question answering. > Learn More	<ul style="list-style-type: none"> > Basic understanding of deep learning concepts. > Basic understanding of language modeling and transformers.
		Tools, Libraries, Frameworks Languages Duration Price Certificate
	NVIDIA NeMo English 6 hours \$30 (excludes tax, if applicable) Yes	
Prompt Engineering With Llama 2	In this course, you'll interact with and prompt engineer Llama 2 models to analyze documents, generate text, and be an AI assistant. > Learn More	Experience with deep learning training using Python.
		Tools, Libraries, Frameworks Languages Duration Price Certificate
	Llama 2, HuggingFace English 3 hours \$30 (excludes tax, if applicable) N/A	

[Back](#)

Course Name	Description	Prerequisites			
Graphics and Simulation					
Assemble a Simple Robot in NVIDIA Isaac Sim™	<p>In this course, you'll step through the "Assemble a Simple Robot" tutorial to rig a two-wheel mobile robot in a live NVIDIA Isaac Sim GPU environment.</p> <p>> Learn More</p>	A Windows or Linux computer with the ability to install Omniverse Launcher and Omniverse applications; internet bandwidth sufficient to support the Isaac Sim client/server stream (performance will vary).			
Course Name	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	NVIDIA Isaac Sim	English	30 minutes	Free	N/A
Build Beautiful, Custom UI for 3D Tools on NVIDIA Omniverse™	<p>Experience the NVIDIA Omniverse development platform for builders and creators of virtual worlds. Become a master in UI with a deep dive into NVIDIA Omniverse Kit's powerful omni.ui suite of tools and frameworks. In this self-paced course, you'll build your own custom UI for workflows in Omniverse with hands-on exercises.</p> <p>> Learn More</p>				Basic familiarity with Python (helpful, not required). Suggested materials to satisfy prerequisites: The Python Tutorial.
Course Name	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	Omniverse Code, Visual Studio Code, Python, and the Python Extension	English, Simplified Chinese	90 minutes	Free	N/A
Develop, Customize, and Publish in NVIDIA Omniverse With Extensions	<p>Want to change the functionality and user interface (UI) of NVIDIA Omniverse? Learn how to customize the Omniverse experience with extensions using Python code.</p> <p>> Learn More</p>				A basic understanding of Python. A basic understanding of computer graphics is useful but not required.
Course Name	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
	Omniverse Code, Visual Studio Code, Python, and the Python Extension	English	8 hours	Free	Yes

Back

Course Name	Description	Prerequisites			
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
Easily Develop Advanced 3D Layout Tools on NVIDIA Omniverse	<p>Get hands-on experience with NVIDIA Omniverse—the platform for connecting and creating physically accurate, 3D virtual worlds. See how easy it is to create your own custom scene layout tools in Omniverse Code with a few lines of Python script. In this self-paced course, you'll build your own custom scene layout in Omniverse with hands-on exercises in Omniverse Code and Python.</p> <p>> Learn More</p>	A basic understanding of computer graphics concepts—such as vertices, meshes, and RGB values—and an understanding of fundamental programming concepts in Python like functions, loops, dictionaries, and arrays.			
Essentials of Developing Omniverse Kit Applications	<p>In this course, participants will learn about kit files and how to create one, how to add extensions to applications, how to define the layout of an application and how to package and distribute an application.'</p> <p>> Learn More</p>	<ul style="list-style-type: none"> > A basic understanding of Python > A basic understanding of computer graphics is useful but not required. > Creating an extension for Omniverse. > Using Github. > How to use terminal commands. 			
Getting Started With USD for Collaborative 3D Workflows	<p>Learn how to generate a scene using human-readable Universal Scene Description ASCII (.USD) files.</p> <p>Upon completion, you'll be able to create your own scenes within the USD framework and will have a strong foundation to use it in applications, such as NVIDIA Omniverse, Maya, Unity, and Unreal Engine.</p> <p>> Learn More</p>	<p>Visual Studio Code and Python</p>	<p>English</p> <p>90 minutes</p>	<p>Free</p>	N/A
How to Build Custom 3D Scene Manipulator Tools on NVIDIA Omniverse	<p>See how you can build advanced tools on the modular, easily extensible Omniverse platform. You'll learn from the Omniverse developer ecosystem team how you can extend and enhance the 3D tools you know and love today. In this self-paced course, you'll build your own custom scene manipulator tools in Omniverse with hands-on exercises writing a few lines of Python code.</p> <p>> Learn More</p>	<p>Universal Scene Description</p>	<p>English, Simplified Chinese</p>	<p>2 hours</p>	N/A
	<p>Omniverse Code, Visual Studio Code, Python, and the Python Extension</p>	<p>English, Simplified Chinese</p>	<p>90 minutes</p>	<p>Free</p>	N/A

Back

Course Name	Description	Prerequisites		
Introduction to Robotic Simulations in NVIDIA Isaac Sim	<p>Robotic automation has enjoyed great success in recent years with increasing hardware capabilities driving innovation in simulation and machine learning. In this course, we introduce you to Isaac Sim, NVIDIA Omniverse's solution for simulation and robotics.</p> <p>You'll learn how to tap into the simulation loop of a 3D engine and initialize experiments with objects, robots, and physics logic. This can be done programmatically using Omniverse Kit and Pixar USD commands, but the course will use Isaac Sim Core to wrap these low-level operations in an object-oriented fashion. By the end of the course, you'll be able to simulate and control NVIDIA JetBot™ and Franka Emika robots and coordinate them together to perform a handoff.</p> <p>The skills covered in this course are direct prerequisites for working with Isaac Gym and create a good starting point for exploring Isaac Sim and other Omniverse applications. The course is great for those interested in 3D scene specification and robotic simulation, but it's also useful for researchers looking to expand their toolkits and seasoned developers interested in exploring design patterns for Omniverse Kit development.</p>	<ul style="list-style-type: none"> > Intermediate knowledge and general comfort with Python 3. This includes familiarity with functions, classes, and basic design patterns. > Comfort with NumPy arrays and basic matrix operations. > A Windows or Linux machine with NVIDIA Omniverse and the Omniverse Streaming Client app. 		
Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
Isaac Sim, Omniverse Kit, NumPy	English, Simplified Chinese	4 hours	\$30 (excludes tax, if applicable)	N/A
Synthetic Data Generation for Training Computer Vision Models	<p>How much data is enough? This is a common question when fine-tuning or training computer vision models. In cases where data collection is a limiting factor, we can use synthetic data! NVIDIA Omniverse Replicator streamlines synthetic data generation (SDG) using 3D assets into a single application, with the ability to modify the appearance and format of the data. This lab highlights one of the ways deep learning tools and Omniverse can be used together to streamline deep learning workloads.</p>	<ul style="list-style-type: none"> > Intermediate understanding of Python (including classes, objects, and decorators): learn about this topic from the Python.org tutorials > Basic understanding of Machine Learning and Deep Learning concepts and pipelines: learn about this topic from the "Deep Learning Demystified" video 		
Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
NVIDIA Omniverse Replicator, NVIDIA Triton Inference Server, PyTorch	English	3 hours	\$30	N/A

Back

Workshop Name	Description	Prerequisites				
Infrastructure						
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
Introduction to AI in the Data Center	<p>Explore AI, GPU computing, NVIDIA AI software architectures, and how to implement and scale AI workloads in the enterprise data center.</p> <p>> Learn More</p>	Artificial intelligence, machine learning, deep learning, GPU hardware and software	English	4 hours	\$49 (excludes tax, if applicable)	Available
Introduction to NVIDIA DOCA™ for DPUs	<p>The NVIDIA DOCA Software Framework lets developers rapidly create applications and services on top of NVIDIA BlueField data processing units (DPUs). Together, DOCA and the BlueField DPU deliver breakthrough networking, security, and storage performance with a comprehensive, open development platform.</p> <p>In this self-paced course, you'll learn the basic concepts of DOCA as a platform for accelerated data center computing on BlueField DPUs. Upon completion, participants will be equipped with introductory knowledge that will enable you to begin using DOCA and DPUs to develop applications that accelerate your data centers services.</p> <p>> Learn More</p>	<ul style="list-style-type: none"> > Familiarity with software architecture and how it relates to and executes on hardware. > Suggested materials to satisfy prerequisite: <ul style="list-style-type: none"> • Enterprise Data Center Networking • Data Center: Overview • Data Center: Virtualization > Some working knowledge of data center networking. > Suggested materials to satisfy prerequisite: <ul style="list-style-type: none"> • Introducing How Computers Work • Hardware Acceleration • Software Execution and Computing 				
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certificate
		NVIDIA DOCA SDK	English, Simplified Chinese	2 hours	Free	N/A

Back

Instructor-Led Workshops for Administrators

Workshop Name	Description	Prerequisites				
AI and Data Science						
NVIDIA AI Enterprise Administration: Public Training		This hands-on training course explores architecture, installation, configuration, operation, and management of NVIDIA AI Enterprise.				
		> Learn More				
Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam		
N/A	English	12 hours	\$1,500	N/A		
Cluster Administration						
Base Command Manager		This course provides an overview of Base Command Manager, including managing nodes and software images, monitoring devices and jobs, managing users, and configuring workload managers.				
		> Learn More				
Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam		
Base Command Manager	English	12 hours	Contact us	N/A		
Ethernet Cumulus						
Cumulus® Linux: Public Bootcamp		Learn how to install, deploy, configure, and troubleshoot Cumulus-based networks. This course offers a perfect blend of hands-on training and theoretical education.				
		> Learn More				
Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam		
Cumulus Linux switches	English	12 hours	\$1,500	Available		
Cumulus Linux: Private Workshop		In this hands-on private training, you'll learn about NVIDIA Cumulus OS architecture, installation, configuration, operation, and management of Cumulus Linux running on NVIDIA switches.				
		> Learn More				
Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam		
Cumulus Linux switches	English	20 hours	Contact us	Available		
NVIDIA Cumulus Linux: Customized Advanced Training		This course focuses on how to build and operate a state-of-the-art data center or storage fabric with emphasis on troubleshooting. The course covers advanced topics such as filtering, quality of service (QoS), Ethernet VPN multihoming (EVPN-MH), monitoring, and active testing.				
		> Learn More				
Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam		
Cumulus Linux switches	English	12 hours	Contact us	N/A		

[Back](#)

Workshop Name	Description	Prerequisites			
InfiniBand					
Workshop Name	Description	Tools, Libraries, Frameworks	Languages	Duration	Price
InfiniBand Customized Course	In this course, you'll learn about InfiniBand architecture and how to manage, monitor, and troubleshoot your InfiniBand network. > Learn More	InfiniBand networks	English	16 hours	Contact us
					Available
Networking Professional Customized Training	In this course, you'll learn about InfiniBand and Cumulus architecture and how to manage, monitor, and troubleshoot triad deployment-based networks. > Learn More	InfiniBand networks	English	16 hours	Contact us
					N/A
NVIDIA DGX					
Workshop Name	Description	Tools, Libraries, Frameworks	Languages	Duration	Price
NVIDIA DGX H100/A100 Administration: Private Workshop	This course provides an overview of the NVIDIA DGX A100 system and NVIDIA DGX Station™ A100, tools for in-band and out-of-band management, NGC, the basics of running workloads, and specific management tools and command-line interface (CLI) commands. In addition, this course includes content on Multi-Instance GPU (MIG), managing storage, performance validation, and other system management tools and concepts. > Learn More	DGX A100 system and DGX Station A100	English	16 hours	Contact us
					N/A
NVIDIA DGX H100/A100 Administration: Public Workshop	This course provides an overview of the DGX A100 system and DGX Station A100's tools for in-band and out-of-band management, the basics of running workloads, specific management tools, and CLI commands. > Learn More	DGX A100 system and DGX Station A100	English	16 hours	\$1,500
					N/A

Back

Workshop Name	Description	Prerequisites				
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
NVIDIA DGX BasePOD Administration: Private Workshop	This course provides an overview of DGX POD components and related processes, including the NVIDIA DGX A100 system, InfiniBand and ethernet networks, tools for in-band and out-of-band management, NGC, the basics of running workloads, and specific management tools and CLI commands. It includes instructions for managing vendor-specific storage per the architecture of your specific POD solution.	DGX POD cluster	English	16 hours	Contact us	N/A
> Learn More						
NVIDIA DGX SuperPOD™ Administration: Private Workshop	This course is designed to help IT professionals successfully administer all aspects of a DGX SuperPOD cluster, including compute, storage, and networking.	DGX SuperPOD cluster	English	16 hours	Contact us	N/A
> Learn More						

Virtualization

Workshop Name	Description	Prerequisites				
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
NVIDIA AI Enterprise Administration: Public Bootcamp	This course covers the platform and solution overview, hardware and software architecture, deployment options, licensing, temporal and spatial GPU partitioning, scaling, comprehensive validation, management, maintenance, monitoring, and troubleshooting.	NVIDIA AI Enterprise	English	12 hours	\$1,500	N/A
> Learn More						

[Back](#)

Online, Self-Paced Courses for Administrators

Course Name	Description	Prerequisites		
AI and Data Science				
Introduction to AI in the Data Center	Explore an introduction to AI, GPU computing, NVIDIA AI software architecture, and how to implement and scale AI workloads in the data center.	None		
	> Learn More			
Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
N/A	English	4 hours	\$49	Available
NVIDIA AI Enterprise Administration	This course covers the platform and solution overview, hardware and software architecture, deployment options, licensing, temporal and spatial GPU partitioning, scaling, comprehensive validation, management, maintenance, monitoring, and troubleshooting.	To gain the most value from this course, the target audience should have a working knowledge in the following domains: <ul style="list-style-type: none"> > Data Center Infrastructure: Servers, Storage, Networking, GPUs, Operating Systems. > Virtualization: VMware vSphere. > Containerization: Docker. 		
	> Learn More			
Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
NVIDIA AI Enterprise	English	8 hours	\$99	N/A
NVIDIA NeMo Framework	Whether you're looking to build LLMs from scratch or to customize community models and deploy them, this course offers the necessary knowledge and skills to accomplish your objectives.	<ul style="list-style-type: none"> > Fundamentals of model parallelism methods for LLMs. > Fundamentals of transformers and conversational AI systems. > Programming languages: Python. > Deep learning frameworks: PyTorch. 		
	> Learn More			
Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
NVIDIA NeMo	English	8 hours	Free	N/A
Cluster Administration				
Base Command Manager	This course is based on NVIDIA Base Command Manager and gives an overview of the cluster management tools, Bright View and cluster management shell (CMSH).	None.		
	> Learn More			
Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
Base Command Manager	English	5 hours	Free	N/A

[Back](#)

Course Name	Description	Prerequisites				
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
Base Command Manager Autoscaling Hybrid Cloud	This course is based on NVIDIA Base Command Manager and gives an overview of extending the cluster to the cloud with Cluster as a service and cluster extension (i.e., hybrid cloud). > Learn More	NVIDIA Base Command Manager	English	3 hours	Free	N/A
Introduction to Base Command Manager	This course is based on NVIDIA Base Command Manager and gives an overview of the usage and components of the software. > Learn More	NVIDIA Base Command Manager	English	3 hours	Free	N/A
DGX						
NVIDIA DGX Cloud	This course is based on NVIDIA DGX Cloud using NVIDIA Base Command Platform. You'll learn to manage users and teams, run single and multi-node jobs, and manage data. > Learn More	DGX Base Command Manage	English	1 hour	Free	N/A
Ethernet						
Linux Networking Fundamentals	Learn the fundamental concepts and commands behind Linux-based open networking. > Learn More	Linux networking concepts	English	6 hours	\$99	N/A

Back

Course Name	Description	Prerequisites				
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
Network Administration With the NVIDIA Onyx™ Switch System	This course provides the required set of skills to configure and manage NVIDIA Ethernet switch systems. You'll learn in depth layer 2 configurations such as VLAN, STP, LAG, and MLAG, as well as how to configure layer 3 features such as BGP.	NVIDIA Onyx	English	3 hours	\$99	N/A
RDMA Over Converged Ethernet (RoCE) From A to Z		Basic understanding of networking concepts and the Open Systems Interconnection (OSI) model.				
		RoCE	English	2 hours	Free	N/A

Graphics and Simulation

NVIDIA Omniverse Enterprise Administration	The course covers the solution overview, hardware and software architecture, deployment options, installation, configuration, licensing, scaling, comprehensive validation, security, management, maintenance, monitoring, and troubleshooting. The instruction and guidance are based on NVIDIA's best practices and cover the critical knowledge and skills for deploying, administering, and managing your Omniverse solution.	None

InfiniBand

InfiniBand Essentials	This self-paced course covers the fundamental first steps into the world of InfiniBand. If you're looking to become more familiar with InfiniBand's benefits, uses, architecture layers, and management concepts, this is the best place to start.	General understanding of networking concepts and principles.

[Back](#)

Workshop Name	Description	Prerequisites				
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
InfiniBand Professional	This course covers the fundamentals of the InfiniBand technology from a usability point of view and builds on the details of the InfiniBand architecture specification. You'll learn how to install, configure, manage, troubleshoot, and monitor your InfiniBand network.					General understanding of networking concepts and principles.

[> Learn More](#)

Management

Data Center Management Made Easy With NVIDIA UFM	Learn about NVIDIA Unified Fabric Manager (UFM) and its capabilities, advantages, and components through a set of interactive learning units, videos, and simulators.	Understanding of InfiniBand fabrics and management concepts			
	> Learn More				
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
	N/A	English	3 hours	\$49	N/A

NVIDIA License System	NVIDIA License System (NLS) is a new licensing solution to support the continued expansion of the NVIDIA enterprise software portfolio. This course will help you to learn about NLS and how you can move from your existing licensing solution to NLS.	> Basic understanding of virtual appliances installation and setup. > Familiarity with web/cloud-based applications. > Familiarity with NVIDIA products like virtual GPU (vGPU) and NVIDIA AI Enterprise.			
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
	Cloud License Service (CLS) and Delegated License Service (DLS)	English	2 hours	Free	N/A

Network

Ansible Essentials for Network Engineers	In this course, you'll explore a variety of Ansible modules and write playbooks specifically adapted to modern data centers. This course includes an exclusive hands-on lab environment and exercises to practice real-world scenarios in real cloud environments.	> Basic Linux administration. > General understanding of networking concepts and principles.			
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
	Ansible	English	3 hours	\$49	N/A

[Back](#)

Workshop Name	Description	Prerequisites				
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
Introduction to Networking	In this course, we'll cover the basics of Ethernet technology and understand how data is forwarded in an Ethernet network. > Learn More			None		
MLXlink and MLXcables Debug Tools	In this course, you'll learn about the MLXlink and MLXcables debug tools. These debug tools are used for both basic link troubleshooting and for analyzing the more complex link characteristics.				Good technical background and understanding of networking hardware.	
		> Learn More				
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam	
	MLXLink and MLXcables	English	2 hours	Free	N/A	
NVIDIA BlueField DPU Administration	Learn the basic concepts of BlueField DPUs as a platform for accelerated data center computing. > Learn More			<ul style="list-style-type: none"> > Basic knowledge and experience in networking concepts and principle. > Basic knowledge and experience in Linux administration. 		
	Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam	
	N/A	English	3 hours	\$49	N/A	

RDMA

The Fundamentals of RDMA Programming	This course allows C programmers to dive into the RDMA programming world without requiring previous experience in networking or RDMA programming. We've also added tips and tricks, as well as do's and don'ts, so the skills you acquire will truly serve you when you need them.	Understanding of C/C++ programming.
	> Learn More	

Back

Certifications

Certification Name	Description	Prerequisites				
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
NVIDIA Certified Associate: AI in the Data Center	This is an entry-level certification that validates foundational concepts of adopting artificial intelligence computing by NVIDIA in a data center environment. The exam is online and remote proctored with 50 questions and a time limit of 60 minutes for completion.	N/A	English	1 hour	\$135	Available
NVIDIA Certified Professional: InfiniBand				A thorough understanding of data center infrastructure and networking.		
		Tools, Libraries, Frameworks	Languages	Duration	Price	Certification Exam
		NVIDIA InfiniBand fabrics	English	1.5 hours	\$220	Available

Ready to Get Started?

To get started with hands-on training, visit
www.nvidia.com/en-us/learn/enterprise

For questions, contact us at nvdi@nvidia.com

© 2023 NVIDIA Corporation and affiliates. All rights reserved. NVIDIA, the NVIDIA logo, BlueField, CUDA, Cumulus, cuOpt, DALI, DGX, DGX POD, DGX Station, DGX SuperPOD, Isaac Sim, Jetson Nano, NEO, NGC, Nsight, NVIDIA DRIVE, NVIDIA DOCA, NVIDIA Merlin, NVIDIA Omniverse, NVIDIA Triton, RAPIDS, TensorRT, and UFM are trademarks and/or registered trademarks of NVIDIA Corporation and its affiliates in the U.S. and other countries. Other company and product names may be trademarks of the respective companies with which they are associated. All other trademarks are property of their respective owners. 3024468.NOV23

